

Stylometric Analysis of Raw Tweets Using Scikit-Learn

William
GA-SEA-DAT2

What is Stylometry?

- **The Problem: Identifying authorship of text through analysis of *writer invariants*.**
- **Two characteristics are necessary for something to qualify as a writer invariant:**
 - It is statistically unique to the same author
 - Remains largely unchanged throughout text
- **Features that may help identify invariants:**
 - **Numeric:** *Average sentence length, syllables, etc*
 - **Synonyms:** *Word choice, use of contractions, etc*
 - **Punctuation:** *Comma, semicolon, or hyphen use*

What is Twitter?

- **Twitter is a social network that is best known for its 140-character limit on user posts.**
- **It does not have any real-name policy, and anyone can easily register more than one account.**
- **In theory, the short character limit on posts should make it harder for writer invariants to surface in new posts.**

The Question(s)

- **In the absence of usernames, what *features* can we best identify tweet authorship with?**
- **Are there features characteristic to tweets that make them easier to attribute?**
- **Hypothesis: Despite the 140-character limit, tweets may have other useful invariants:**
 - ***Frequency of link usage, which domains***
 - ***If added text is placed before/after links***
 - ***Hashtag usage, response to other users***

Some Limitations

- **Prospective user request: “I don't trust anything I can't understand myself.”**
- **Therefore: Simplicity > Accuracy**
- **Feature Combinations: A model with too many features early on makes it hard to tell what the effect of each is.**
- **Realism: A typical stylometry scenario will generally be between suspect and unattributed text passages.**

Procedural Summary

- 1. Data Acquisition:** Acquire tweets from two or more feeds using Twitter's API
- 2. Pre-Processing:** Extract numeric data about possible features discussed earlier
- 3. Model Development:** Create a logistic regression model that can attribute tweets
- 4. Analysis/Iteration:** Attempt to find the *simplest* model that can achieve stated goal
- 5. Result Visualization:** Produce visuals that can convey to possible users how this works

The Data

- **Source:** Via Twitter's API I downloaded two feeds that discuss the same subject matter. This eliminates the temptation to use topic modeling.
- **What Code:** To acquire the tweets in CSV form, I used a python library called “Tweepy” which requires that you have a Twitter account to work with it's API.



Ninja Economics

@NinjaEconomics

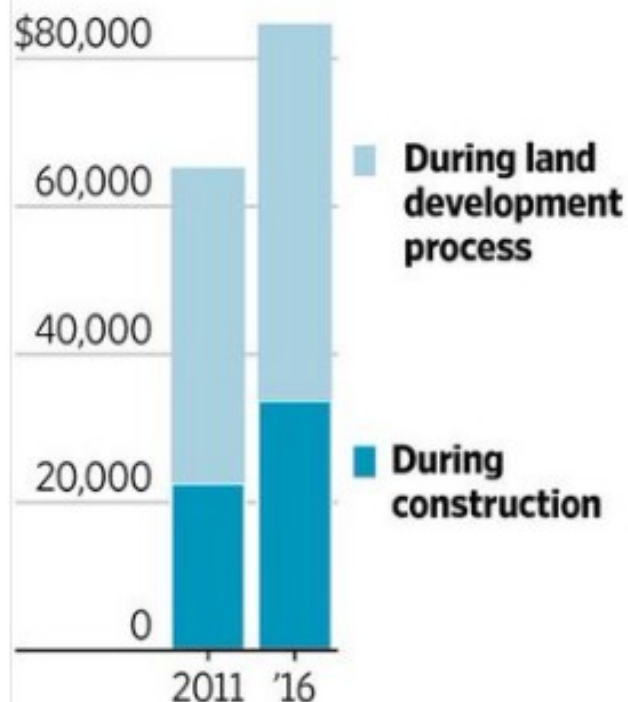
Follow

The average cost for new-home builders to comply with regulations has increased by nearly 30% over the last 5 years.


Cost of Doing Business

Home builders' costs of complying with regulations have jumped over five years, rising at roughly the same rate as the price of a new home.

Regulatory costs of average new single-family home



Growth in costs compared to rise in average home price





Daniel Lin

@DLin71

Follow

There's a San Francisco edition of Monopoly.
You're not allowed to build any homes, and the
game ends when everyone moves to Oakland.

RETWEETS

1,097

LIKES

1,830



Daniel Lin

@DLin71

Follow

What we know about Paris terrorists

- Not Syrian
- Not refugees
- No encryption

What the US is focusing on

- Syrians
- Refugees
- Encryption

RETWEETS

25,517

LIKES

17,535



```

15 def get_all_tweets(screen_name):
16     #Twitter only allows access to a users most recent 3240 tweets with this method
17
18     #authorize twitter, initialize tweepy
19     auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
20     auth.set_access_token(access_key, access_secret)
21     api = tweepy.API(auth)
22
23     #initialize a list to hold all the tweepy Tweets
24     alltweets = []
25
26     #make initial request for most recent tweets (200 is the maximum allowed count)
27     new_tweets = api.user_timeline(screen_name = screen_name, count=200)
28
29     #save most recent tweets
30     alltweets.extend(new_tweets)
31
32     #save the id of the oldest tweet less one
33     oldest = alltweets[-1].id - 1
34
35     #keep grabbing tweets until there are no tweets left to grab
36     while len(new_tweets) > 0:
37         print "getting tweets before tweet ID %s" % (oldest)
38
39         #all subsequent requests use the max_id param to prevent duplicates
40         new_tweets = api.user_timeline(screen_name = screen_name, count=200, max_id=oldest)
41
42         #save most recent tweets
43         alltweets.extend(new_tweets)
44
45         #update the id of the oldest tweet less one
46         oldest = alltweets[-1].id - 1
47
48         print "...%s tweets downloaded so far" % (len(alltweets))
49
50     #transform the tweepy tweets into a 2D array that will populate the csv
51     outtweets = [[tweet.created_at, tweet.text.encode("utf-8")] for tweet in alltweets]
52
53     #write the csv
54     with open('%s_tweets.csv' % screen_name, 'wb') as f:
55         writer = csv.writer(f)
56         writer.writerow(["Time and Date Tweeted", "Raw Tweet Content"])
57         writer.writerows(outtweets)
58
59     pass

```


The Modeling

- **Logistic Regression:** Data is not numerically continuous, quantities and yes/no for features.
- **First:** Analyze twitter feed by chosen features.
- **Next:** Choose which features best distinguish one feed from another, much iteration...
- **Final:** Gather a subset of new tweets from previously selected feeds (ten or so), remove the user handles, and use model to attribute.
- **Iris Dataset:** Very similar problem, in the sense that the aim is to categorize by chosen features.

Example: Syllable Count

```
794
795 word = raw_input('Enter phrase: ')
796 word = word[0].upper() + word[1:].lower()
797 print
798 print word
799
800 # Count the syllables in the word.
801 syllables = 0
802 for i in range(len(word)) :
803
804     # If the first letter in the word is a vowel then it is a syllable.
805     if i == 0 and word[i] in "aeiouy" :
806         syllables = syllables + 1
807
808     # Else if previous letter isn't a vowel
809     elif word[i - 1] not in "aeiouy" :
810
811         # If not the last letter and is a vowel
812         if i < len(word) - 1 and word[i] in "aeiouy" :
813             syllables = syllables + 1
814
815         # Else if it is the last letter and it is a vowel that is not e.
816         elif i == len(word) - 1 and word[i] in "aeiouy" :
817             syllables = syllables + 1
818
819     # Adjust syllables from 0 to 1.
820     if len(word) > 0 and syllables == 0 :
821         syllables == 0
822         syllables = 1
823
824
825 # Display the result.
826 print
827 print "The word contains", syllables, "syllable(s)"
828
```